

CONTROL AUTHORITY BOUNDARIES: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study shared autonomy and robot control. The motivating bet is: Model where authority should shift between policy, controller, and human.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For shared autonomy and robot control, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026; 2020a; 2021a; 2023; 2020b; 2021b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.530	0.470
observed_only_wam	hidden_mode_shift	0.152	0.848
augmentation_as_data_wam	nominal	0.600	0.400
augmentation_as_data_wam	hidden_mode_shift	0.275	0.725
scalar_uncertainty_wam	nominal	0.685	0.315
scalar_uncertainty_wam	hidden_mode_shift	0.367	0.633
proposed_branch_mechanism	nominal	0.785	0.215
proposed_branch_mechanism	hidden_mode_shift	0.585	0.415

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Control Authority Boundaries argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Adaptive impedance control for the haptic shared driving task based on nonlinear mpc, 2020a. <https://doi.org/10.1115/dscc2020-3296>.
- Haptic shared control for human-robot collaboration: A game-theoretical approach, 2020b. <https://doi.org/10.1016/j.ifacol.2020.12.2751>.
- Bayesian controller fusion: Leveraging control priors in deep reinforcement learning for robotics, 2021a. <http://arxiv.org/abs/2107.09822v3>.
- A barrier pair method for safe human-robot shared autonomy, 2021b. <http://arxiv.org/abs/2112.00279v1>.
- A shared control approach based on first-order dynamical systems and closed-loop variable stiffness control, 2023. <https://doi.org/10.1007/s10846-023-02023-w>.
- Safety-critical control for human-robot interaction in shared workspaces using control barrier functions, 2026. <https://doi.org/10.66238/ijcbs40>.